Athlete Peak Performance: An Analysis of Player Prime Years across Sports

NAU NORTHERN ARIZONA UNIVERSITY
Department of Mathematics and Statistics

Jacob Berlin, Eli Vatsaas

The Department of Mathematics and Statistics, Northern Arizona University, Flagstaff , AZ

NAU NORTHERN ARIZONA UNIVERSITY
Office of Undergraduate Research and Creative Activity

Abstract

The peak performance period of an athlete—commonly known as their "prime"—represents one of the most captivating aspects of sports fandom. This study investigates the relationship between age and athletic performance across different sports leagues, focusing not on biological physical capacity but on actual competitive production within each sport's context. Drawing from "SportsDataverse" R packages and supplementary sources like wikidataR, we constructed a comprehensive dataset featuring 27,687 athletes competing across 9 leagues in 6 sports. For each athlete, we documented age and league-specific performance metrics at the seasonal level, enabling precise analysis of performance trajectories. Our analysis reveals distinct patterns in the timing and duration of athletic primes across different leagues, highlighting both universal trends and league-specific variations. These findings provide valuable insights for talent development strategies, contract valuations, performance expectations, and fan engagement. Understanding the typical prime trajectory for athletes in various sports has practical applications for teams, analysts, and the broader sports community in predicting and contextualizing athletic performance.

Introduction

One of the most intriguing aspects of professional sports are primes. An individual athlete that is performing at their absolute best can sometimes singlehandedly change the outcome of matches, making primes invaluable. As fans of sports, we were interested in whether such a critical piece to a competitor's career varied by sport. With this study, we aimed to answer the following questions:

- Do athletes of different sports enter their prime at different ages?
- Is there a significant difference in when athletes of different sports reach their **absolute peak**?
- How different is the lifespan of athletes' primes among different sports?

Defining an Athlete's Prime

• **Peak Season**: Centered around a player's best season based on a key performance metric.

• **70% Threshold**: Seasons are included in the prime if performance remains ≥ 70% of peak.

• **Bidirectional Expansion**: We extend the prime forward and backward from the best season as long as performance stays above the threshold.

• **Skip Rule** (excluding chess): Athletes are allowed one "skip" season in either direction (e.g. due to injury) where performance dips below 70% but rebounds the following year.

Methods

Data Preparation:

For every professional league we analyzed, we used separate R package to obtain and clean data from each league individually. The general steps for preparing each data set for analysis included:

1. Creating a unique ID for each athlete
2. Compiling each athlete's "per game" stats into "per season" stats
3. Ensuring each athlete had ages associated with each season
    - If an athlete did not have age/birthdate information we used wikidataR to obtain birthdate
4. Calculate each athlete's respective sport's performance statistic
5. Filtered out players that didn't make their debut by 23 and all seasons under 18 (resulted in 6045 athletes)

After we completed this process for each league, we combined them all into one large data set.

Spline Fitting:

The Problem:

Player performance data is inherently volatile, with significant season-to-season fluctuations due to injuries, team changes, and natural variance. We knew we needed a method to capture the underlying career trends while filtering out this noise. Smoothing splines were an obvious choice for this.

Our first attempt relied solely on Generalized Cross-Validation (GCV) to fit player career trajectories. This created two significant issues:

1) Too Many Straight Lines: GCV often produced overly simplified linear fits that failed to capture the natural rise and fall of player careers, which is necessary for our analysis

2) Artificial Entry Peaks: For many players, GCV models incorrectly predicted peak performance at their entry year, contradicting their actual career patterns

These issues required developing a more sophisticated approach that could identify and properly handle different career patterns while preventing unrealistic career paths.

Sharp Peak Detection:

We identify players with exceptional performance spikes using a peak prominence measure:

$$\text{Peak Prominence} = (\text{peak\_value} - \text{average\_non\_peak}) / (\text{max\_value} - \text{min\_value})$$

Example: A basketball player with career values ranging from 10 to 50 points per game (range = 40), with a peak of 50 and average performance of 20 in all other seasons.

• Peak Prominence = (50 - 20) / 40 = 30 / 40 = 0.75
• Interpretation: The gap between their peak and typical performance represents 75% of their entire career performance range

We classify careers with prominence ≥ 0.7 as "sharp peak" careers requiring specialized handling.

Two-Path Fitting Strategy:

Path A: Sharp Peak Careers

For players with identified sharp peak seasons:

1. Apply weighted spline fitting with emphasis on peak seasons
2. Give 5× weight to the peak year and adjacent seasons
3. Test degrees of freedom from 3-8 and select model with lowest weighted error
4. Method: Sharp Peak
    - This attempts to balance peak preservation with overall career trajectory smoothness

Path B: Standard Careers

For players with more consistent career trajectories:

1. Try Generalized Cross-Validation (GCV) fitting first
    - Method: Balanced GCV
    - This attempts to find an optimal balance between fit and smoothness using statistical principles
2. If that fails validation, try lower smoothing parameter (spar=0.3)
    - Method: Wiggly Low Spar
    - This attempts to create a more flexible curve that can better capture fluctuations
3. If both fail, test all degrees of freedom (3-8) and select model with lowest error
    - Method: Fallback
    - This attempts to systematically explore all complexity options to find the best possible fit
4. If no models pass both validation rules:
    - We select the best model that passed complexity constraints but failed the entry peak rule

This dual path approach attempts to apply the most appropriate spline based on each player's unique career pattern. Resulting in consistent career curves that accurately capture player development patterns across all sports and positions, despite their vastly different career paths.

League Distribution of Splining Methods

Logo	League	Total Players	Sharp Peak	Balanced GCV	Wiggly Low Spar	Fallback
			Spline Fit Methods (Count & Percentage)			
	CHESS - Male	185	7 (3.8%)	109 (58.9%)	56 (30.3%)	13 (7.0%)
	CHESS - Female	65	6 (9.2%)	32 (49.2%)	26 (40.0%)	1 (1.5%)
	NWSL	112	24 (21.4%)	42 (37.5%)	43 (38.4%)	3 (2.7%)
	MLS	274	60 (21.9%)	112 (40.9%)	96 (35.0%)	6 (2.2%)
	NBA	558	153 (27.4%)	246 (44.1%)	112 (20.1%)	47 (8.4%)
	NFL	2471	768 (31.1%)	1021 (41.3%)	571 (23.1%)	111 (4.5%)
	NHL	649	205 (31.6%)	221 (34.1%)	175 (27.0%)	48 (7.4%)
	WNBA	215	76 (35.3%)	63 (29.3%)	52 (24.2%)	24 (11.2%)
	MLB	1516	635 (41.9%)	434 (28.6%)	305 (20.1%)	142 (9.4%)
	Overall	**6045**	**1934 (32.0%)**	**2280 (37.7%)**	**1436 (23.8%)**	**395 (6.5%)**

Table 1: This table shows the number of players in each league and number of players with each spline method. The numbers in parentheses are the percentages of players fit with the given method relative to the total number of players in that league. This gives context to the player career paths, and thus how our splining methods differ between leagues.

Sample of Identified Primes by Actual and Spline Predicted Career Trajectories

Player Spline Method Indicated

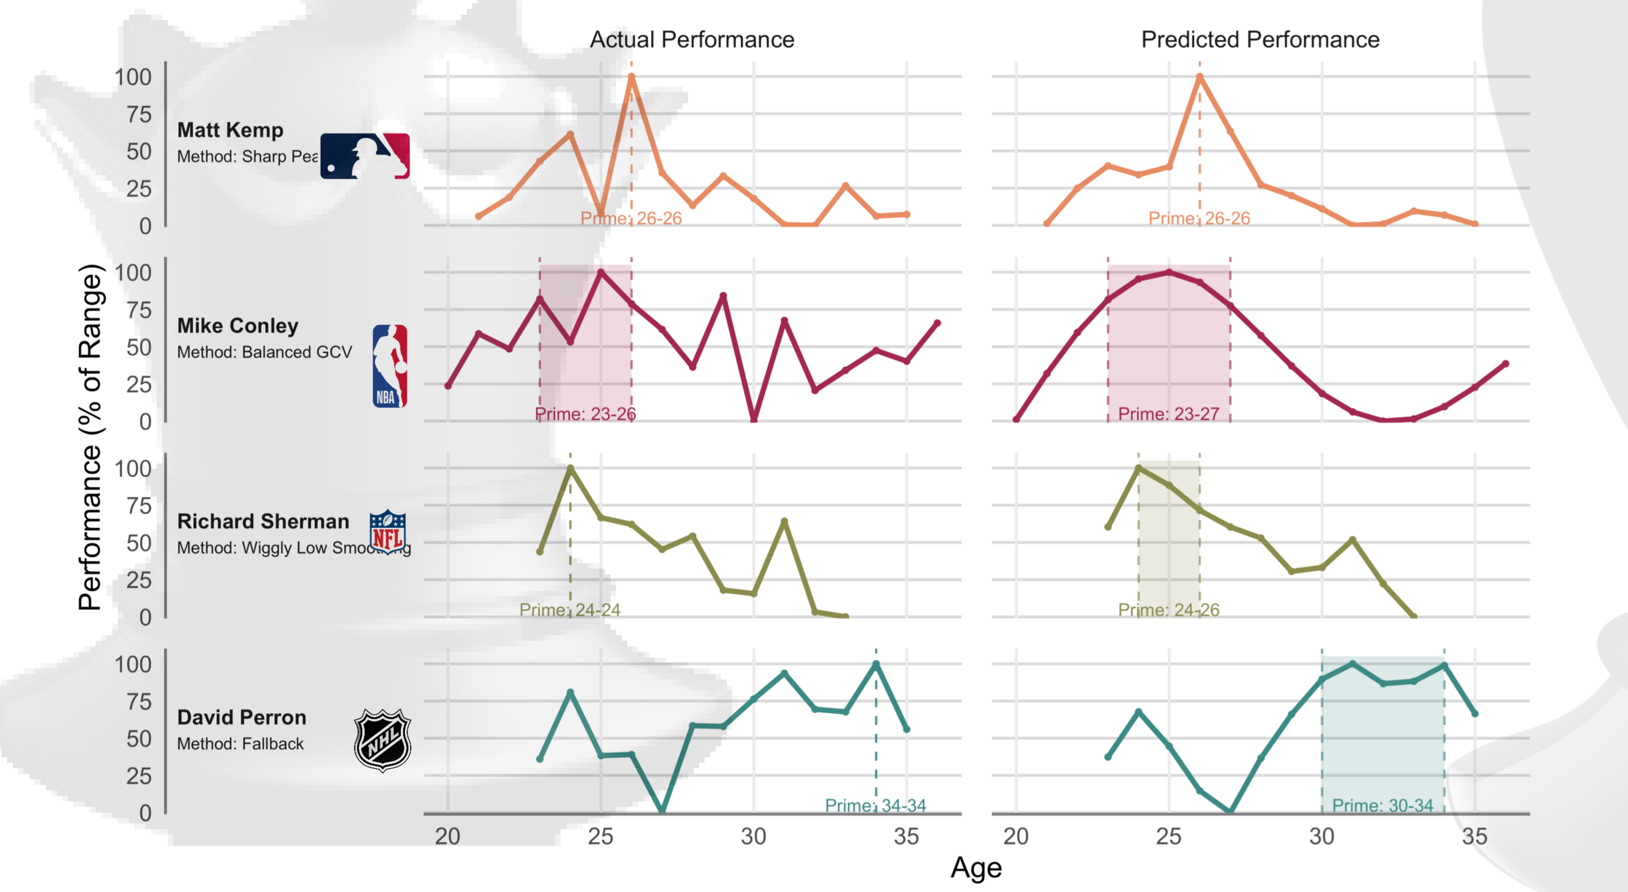

Figure 1: Career trajectories for 4 selected players showing how different spline methods identify prime years. The visualization contrasts actual performance with natural variance (left) against our fitted smoothing splines (right) that reveal underlying career patterns. Each athlete demonstrates a distinct career trajectory handled by distinct splining methods: Sharp Peak for Matt Kemp's dramatic single-season spike; Balanced GCV for Mike Conley's simple career path; Wiggly Low Smoothing for Richard Sherman's volatile performance pattern; and a Fallback method for David Perron's hard to define trajectory. The shaded regions highlight identified prime windows, with spline-based predictions (right) working through the noise in the raw data. This approach demonstrates why statistical smoothing is essential for accurately defining player primes—it filters out seasonal noise, injuries, and team context to capture true performance trends, offering more reliable frameworks for the analysis.

Statistical Testing of Peak Age Differences

The Problem

We wanted to determine whether the observed differences in peak performance ages across sports leagues were meaningful or simply due to random chance. Could we confidently state that some leagues truly peak earlier or later than others?

To answer this question, we examined our data and found two important challenges:
1. Unequal Variability: Some leagues showed much more variability in peak ages than others. For example, chess players' peak ages varied much more widely than the other leagues' peak ages.

2. Uneven Sample Sizes: Our data included very different numbers of athletes across leagues (ranging from just 65 CHESS - Female players to over 2,400 NFL players).
These challenges meant that standard statistical tests would be unreliable and potentially misleading.

Our Approach
To solve this, we implemented the following method that accommodates both uneven sample sizes and unequal variability between groups:

- Built a linear model that predicted peak age based on league
- Applied heteroscedasticity-consistent standard errors (specifically the HC3 variant), which adjusts the uncertainty in our estimates to account for different variance across leagues
- Used estimated marginal means to handle the unbalanced sample sizes, ensuring each league's average peak age was properly estimated regardless of how many athletes it contained
- Conducted pairwise comparisons between every pair of leagues to identify which differences were statistically significant

This approach allowed us to make reliable comparisons between leagues despite dramatically different sample sizes and variability patterns, ensuring our conclusions about peak-age differences were statistically sound.
By accounting for these challenges, we could confidently determine which leagues truly differed in their peak performance ages, without being misled by the uneven nature of the data.

Career Progression Timeline by League

Average age at key career milestones

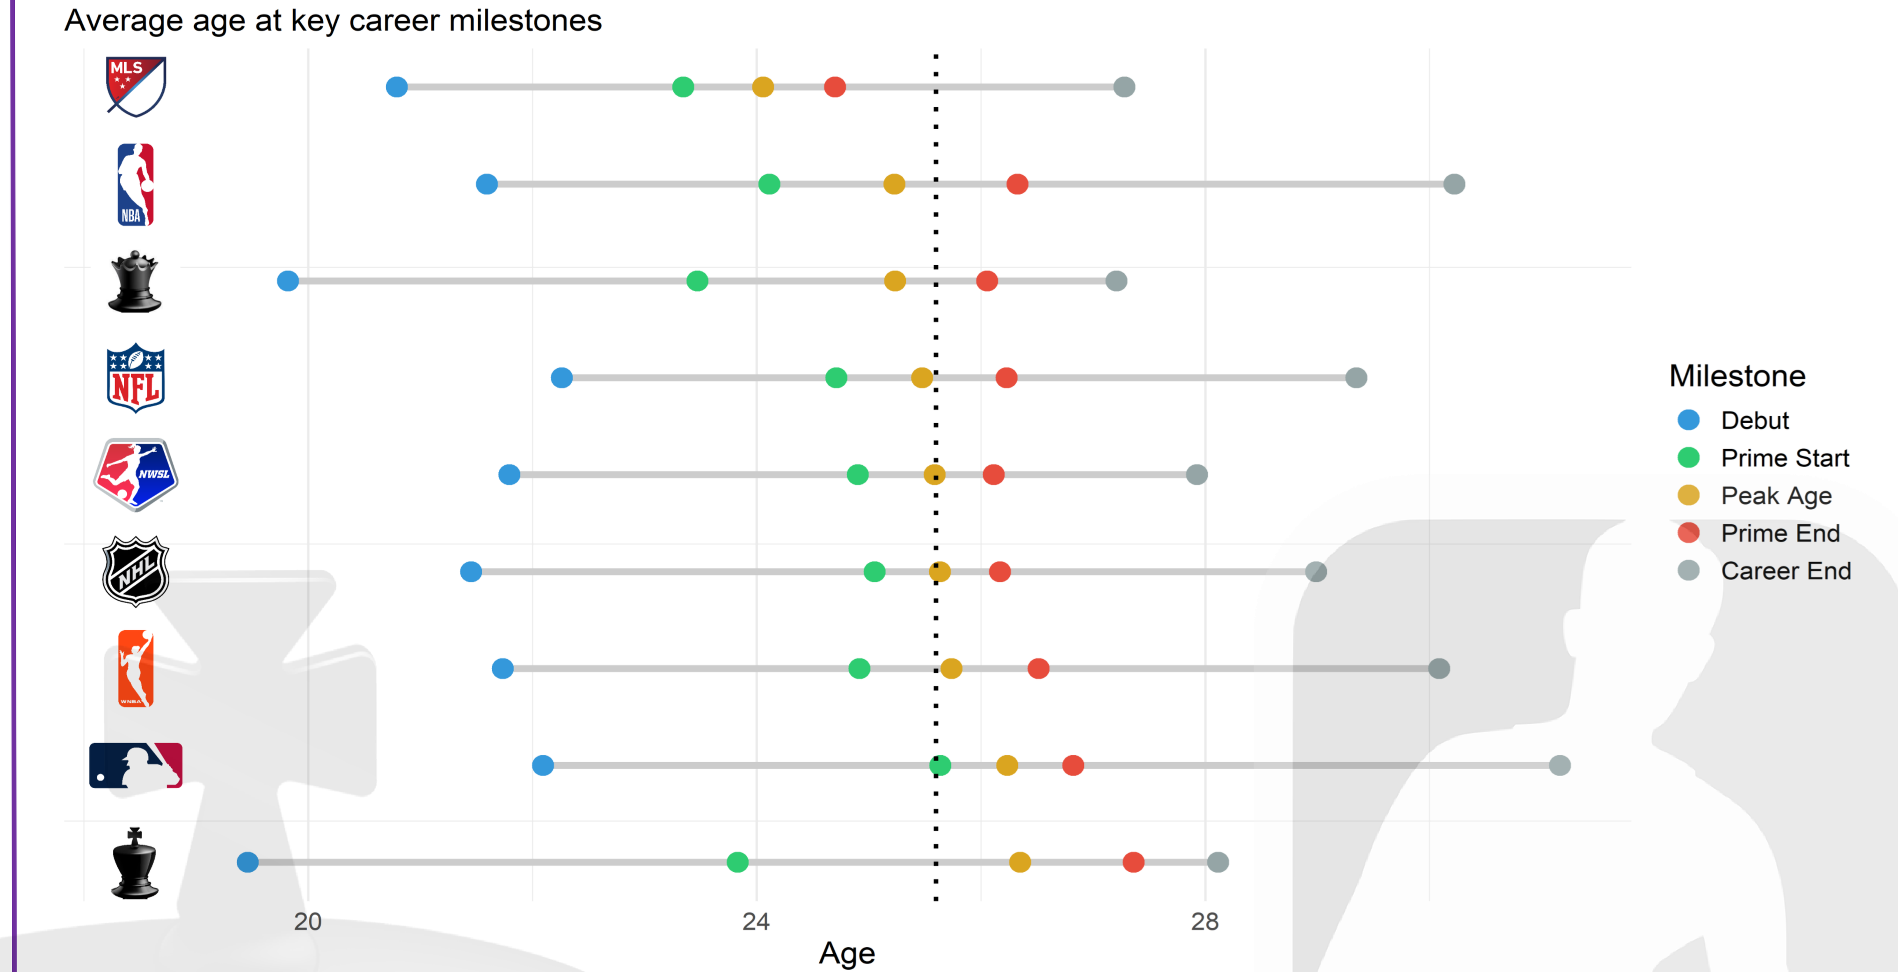

Figure 2. This timeline visualization shows the average age progression through major career milestones for athletes across different leagues. Each row represents a league, with colored points marking the average age at debut, start of prime, end of prime, and career end. The dotted line denotes the average peak age among all sports analyzed. This visualization provides insight into the typical career trajectory and timing across different leagues. We can see that MLS athletes reach their peak at a much better at a younger age than the other leagues, which is reasonable considering their young debut age. Chess players also stand out due to not only their young debut age, but their lengthy prime which results from a reduced performance drop-off in mental sports compared to physical sports.

Conclusion

Our analysis reveals significant and meaningful differences in the age at which athletes reach their peak performance across sports leagues. Peak performance ages vary substantially, with statistically significant differences between one-third of all possible league pairs. These differences hold even after accounting for the statistical challenges of unequal variances and large disparities in sample sizes between leagues.

Through our statistical approach, we can confidently distinguish true differences in peak performance timing from random variation. The patterns we've identified help quantify what many intuitively understand—that different sports demand different physical and technical attributes which peak at different ages.

Understanding these sport-specific prime trajectories provides valuable context for team management strategies around talent development and contract planning, while offering athletes more information for career planning. This research shows a data-driven approach to understand the relationship between age and elite athletic performance, illuminating how career arcs truly differ across competitive contexts.

Significant Differences in Peak Performance Age Across Leagues

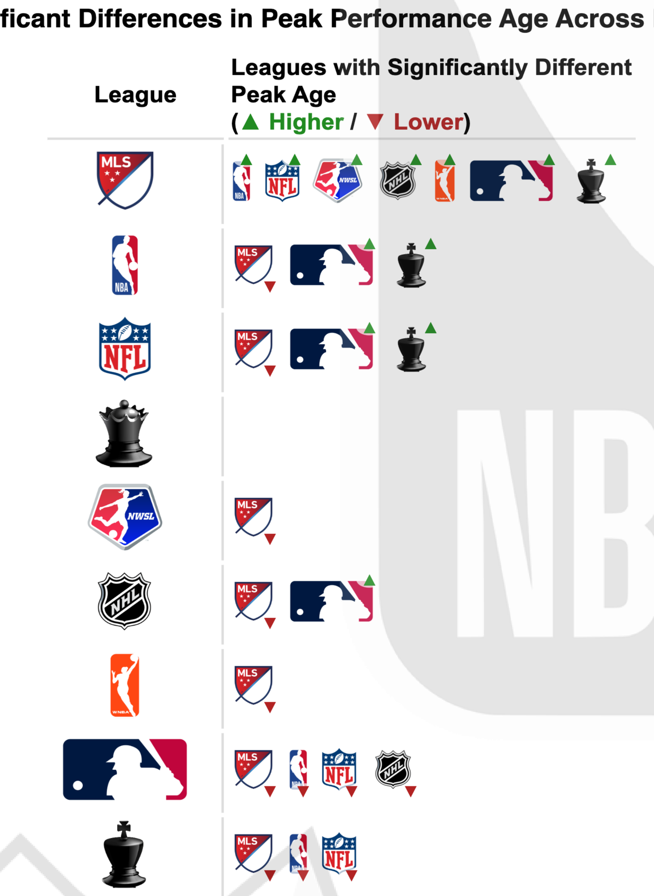

Table 2: Comparison of mean peak performance age. The 'League' column lists leagues ordered lowest-to-highest by mean. The second column identifies leagues with a significantly different mean peak performance age ($p < 0.05$) compared to the respective 'League' in the row. Arrows on these logos indicate the direction of the difference: ▲ means significantly higher age, ▼ means significantly lower age.

Acknowledgements

- Thank you to the NAU Department of Mathematics and Statistics
- Thank you to the Office of Undergraduate Research and Creative Activity for giving us the opportunity to display our data analysis
- Thank you to Dr. Robert Buscaglia for guiding us through this project

References

SportsDataverse. (2024). *SportsDataverse GitHub organization* [R packages and data tools]. GitHub. https://github.com/sportsdataverse
(Packages used: baseballR, hoopR, wehoop, hockeyR)

Chess.com. (2024). *Published Data API*. https://support.chess.com/en/articles/9650547-published-data-api

Richardett, T., & American Soccer Analysis. (2025). *itscalledsoccer* (Version 0.3.1) [R package]. CRAN. https://cran.r-project.org/package=itscalledsoccer

Shafee, T., Keyes, O., Signorelli, S., Lum, A., Graul, C., & Popov, M. (2021). *WikidataR* (Version 2.3.3) [R package]. CRAN. https://cran.r-project.org/package=WikidataR

Carl, S., & Baldwin, B. (2024). *nflfastR: Functions to Efficiently Access NFL Play by Play Data* (Version 5.0.0) [R package]. CRAN. https://cran.r-project.org/package=nflfastR

Extra Info

Even more interested in peak performance among sports? Don't agree with our 70% threshold? Want to look up your favorite player? Scan the QR code to visit our online dashboard where you can explore more about how primes differentiate among sports as well as individual players!

https://elivatsaas.github.io/peakPerformR-dashboard/